

The “No Stone Unturned” 17-Month Reinforcement Learning Protocol

From Real Numbers to Distributed Multi-Agent Systems

A Comprehensive Mathematical & Practical Journey
Implementation × Optimization × Analysis

February 12, 2026

“You asked for it. We are doing this.”

THE COMMITMENT

- **Duration:** 17 Months (68 weeks)
- **Daily Investment:** 2-3 hours
- **Three Threads:** Implementation (The Body), Optimization (The Muscle), Analysis (The Soul)
- **Outcome:** From “Student” to “Specialist”

Contents

1	Introduction: The Philosophy	4
1.1	The Three-Thread Approach	4
1.2	The “Glue” Philosophy	4
1.3	Time Management Protocol	4
2	Phase 1: The Axioms & The Grid	5
2.1	Weeks 1-2: Foundations	5
2.2	Weeks 3-4: The Bellman Equation	5
2.3	Weeks 5-6: The Loop (MDPs)	6
2.4	Weeks 7-8: First Algorithms	6
2.5	Weeks 9-10: Monte Carlo & TD	7
2.6	Weeks 11-12: Benchmarking	7
3	Phase 2: The Gradient & The Manifold	8
3.1	Weeks 13-14: PyTorch from Scratch	8
3.2	Weeks 15-16: DQN (Deep Q-Network)	8
3.3	Weeks 17-18: DQN Stabilization	8
3.4	Weeks 19-20: Policy Gradients (PG)	9
3.5	Weeks 21-22: Actor-Critic (A2C)	9
3.6	Weeks 23-24: Review & Catch-up	10
4	Phase 3: The Trust Region & The Hierarchy	11
4.1	Weeks 25-26: TRPO (Theory)	11
4.2	Weeks 27-28: Conjugate Gradient (CG)	11
4.3	Weeks 29-30: PPO (Implementation)	11
4.4	Weeks 31-32: Hierarchical RL (HRL)	12
4.5	Weeks 33-34: Implicit Differentiation	12
4.6	Weeks 35-36: The Probabilistic Wall	13
5	Phase 3.5: The “No Reward” World	14
5.1	Weeks 33-34: Goal-Conditioned RL	14
5.2	Weeks 35-36: Contrastive RL	14
5.3	Weeks 37-38: Unsupervised Skill Discovery	15
6	Phase 3.6: The “Offline” World	16
6.1	Weeks 39-40: Offline RL Basics	16
6.2	Weeks 41-42: Conservative Q-Learning	16
6.3	Weeks 43-44: Implicit Q-Learning	16
7	Phase 4: The Continuous World & The “Walls”	18
7.1	Weeks 45-46: Control as Inference	18
7.2	Weeks 37-38: Continuous Control (SAC)	18
7.3	Weeks 39-40: The Reparameterization Trick	19
7.4	Weeks 41-42: The Control Theory Wall	19
7.5	Weeks 47-48: LQR (Linear Quadratic Regulator)	20
7.6	Weeks 43-44: Safety & Lyapunov	20

7.7	Weeks 49-50: Model-Based RL (Dreamer)	20
7.8	Weeks 45-46: The Game Theory Wall	21
7.9	Weeks 47-48: The Minimax Theorem	21
8	Phase 5: The Grand Finale	22
8.1	Weeks 49-50: The Systems Wall	22
8.2	Weeks 51-52: V-Trace Correction	22
8.3	Weeks 51-52: Multi-Agent (MARL)	22
8.4	Weeks 53-56: Distributed Systems	23
8.5	Weeks 53-56: CAPSTONE PROJECT	24
8.6	Weeks 57-64: Review & Write-up	24
9	The Complete Algorithm Arsenal	25
10	The Reading Stack	26
10.1	Implementation (The Body)	26
10.2	Optimization (The Muscle)	26
10.3	Analysis (The Soul)	26
10.4	Specialized Topics	26
11	Critical Survival Rules	27
11.1	The 20-Minute Rule	27
11.2	The “Steal the Tools” Mindset	27
11.3	Code “From Scratch” vs. “From Memory”	27
11.4	The Emergency Eject Priority	27
12	Weekly Schedule Template	28
12.1	Day A: The Engineer (Mon, Wed, Fri)	28
12.2	Day B: The Mathematician (Tue, Thu, Sat)	28
12.3	Sunday: The Rest (or The Spillover)	28
13	The Final Warning	29
14	Success Metrics	30
14.1	Month 6 Checkpoint	30
14.2	Month 12 Checkpoint	30
14.3	Month 17 Final Assessment	30
14.4	The “Proof” Sentence	30
15	Conclusion: The Deal	31

1 Introduction: The Philosophy

This is not a casual learning plan. This is a **surgical protocol** designed to take you from understanding “what is a supremum” to proving convergence guarantees for distributed multi-agent reinforcement learning systems.

1.1 The Three-Thread Approach

Every concept in this curriculum is approached from three angles:

1. **Implementation (The Body):** Writing actual code that works
2. **Optimization (The Muscle):** Understanding gradient dynamics, convergence, and stability
3. **Analysis (The Soul):** Rigorous mathematical foundations from Real Analysis, Functional Analysis, and Measure Theory

1.2 The “Glue” Philosophy

Each bi-weekly block includes a “**Glue**” section that explicitly connects all three threads. For example:

“You can’t calculate $\sup S$ (Supremum) if your floating point math overflows (Stability).”

This is not academic; this is survival in production ML systems.

1.3 Time Management Protocol

The A/B Split Strategy:

- **Day A (Mon, Wed, Fri):** The Engineer — Implementation & Optimization
- **Day B (Tue, Thu, Sat):** The Mathematician — Analysis & Theory
- **Sunday:** Rest or spillover

The 20-Minute Rule: If stuck on a mathematical concept for >20 minutes, stop. Google the intuition. Accept the black box. Move forward.

2 Phase 1: The Axioms & The Grid

Months 1-3 | Theme: Discrete Math, Determinism, Perfect Information

Goal: Build a tabular agent while understanding the exact topology of the space it lives in.

2.1 Weeks 1-2: Foundations

Thread	Topic	Tasks	Why?
Implementation	NumPy & Vectorization	Matrix multiplication without loops; Visual Intro to NumPy	Foundation for all array operations
Optimization	Numerical Stability	Overflow/underflow; Demonstrate Soft-max failure	Production code must handle edge cases
Analysis	The Real Numbers	Axiom of Completeness; Supremum vs Max	Mathematical rigor begins here
highlightThe Glue: You can't calculate $\sup S$ if your floating-point math overflows. Stability and mathematical completeness are inseparable.			

2.2 Weeks 3-4: The Bellman Equation

Thread	Topic	Tasks	Why?
Implementation	The Bellman Equation	Solve 4×4 Grid-world; Eq: $V(s) = \max_a \sum_{s'} P(s' s, a) [R + \gamma V(s')]$	Core RL equation
Optimization	Conditioning	Condition Number $\kappa = \frac{\lambda_{\max}}{\lambda_{\min}}$; Valley Geometry	Understanding ill-conditioned systems
Analysis	Sequences	Monotone Convergence Theorem; Bounded increasing sequence converges	Proof foundations
highlightThe Glue: Value Iteration is just a sequence $\{V_0, V_1, V_2, \dots\}$. Monotone convergence guarantees it stops.			

2.3 Weeks 5-6: The Loop (MDPs)

Thread	Topic	Tasks	Why?
Implementation	The Loop (MDPs)	Build Environment API; Transition Matrices $P(s' s, a)$	Standard RL interface
Optimization	Loss Landscapes	Strong Convexity/Smoothness; Jain & Kar Sec 2.4.2	Understanding optimization geometry
Analysis	Metric Spaces	Cauchy Sequences; Completeness of Metric Spaces	Distance and convergence
highlightThe Glue: We define the “distance” between two policies using the Sup-norm metric: $d(V_1, V_2) = \sup_s V_1(s) - V_2(s) $			

2.4 Weeks 7-8: First Algorithms

Thread	Topic	Tasks	Why?
Implementation	First Algorithms	Value Iteration; Verify against theory	Practical implementation
Optimization	Fixed Points	Iterative Solvers; Linear Algebra basics	Solving $x = f(x)$
Analysis	Contraction Mapping	Banach Fixed Point Theorem; Prove Bellman is a Contraction	THE fundamental theorem
highlightThe Glue: CRUCIAL: You will mathematically prove your code <i>must</i> work because $\gamma < 1$ squeezes the metric space.			

Important

The Banach Fixed Point Theorem is the cornerstone of all RL theory. If you understand one theorem deeply, make it this one.

Theorem: Let (X, d) be a complete metric space and $T : X \rightarrow X$ be a contraction mapping (i.e., $\exists \gamma < 1$ such that $d(Tx, Ty) \leq \gamma \cdot d(x, y)$ for all $x, y \in X$). Then:

1. T has a unique fixed point $x^* \in X$
2. For any $x_0 \in X$, the sequence $x_{n+1} = T(x_n)$ converges to x^*
3. The error decreases exponentially: $d(x_n, x^*) \leq \gamma^n d(x_0, x^*)$

2.5 Weeks 9-10: Monte Carlo & TD

Thread	Topic	Tasks	Why?
Implementation	Monte Carlo & TD	“The Cliff Walker”; SARSA vs Q-Learning	Sampling-based methods
Optimization	Noisy Optimization	Intro to Stochastic Approximation; Expectation vs Sample	Handling randomness
Analysis	Topology of $\mathbb{R}^n$	Compactness (Heine-Borel); Closed and Bounded sets	Preventing explosions
The Glue: Q-learning samples are noisy. Compactness ensures our parameters don’t explode to infinity.			

2.6 Weeks 11-12: Benchmarking

Thread	Topic	Tasks	Why?
Implementation	Benchmarking	Stable Baselines 3 comparison; Vectorized Envs	Production standards
Optimization	Floating Point 2.0	Log-Sum-Exp Trick; Numerical stability in RL	Essential for production
Analysis	Uniform Continuity	δ depends only on ϵ ; Stability in discretization	Continuous approximations
The Glue: Production code requires the Log-Sum-Exp trick, or your probabilities will turn to NaN.			

Note

Log-Sum-Exp Trick: Instead of computing $\log(\sum_i e^{x_i})$ directly (which overflows), compute:

$$\text{LSE}(x_1, \dots, x_n) = c + \log \left(\sum_i e^{x_i - c} \right)$$

where $c = \max_i x_i$. This is numerically stable and essential for softmax operations.

3 Phase 2: The Gradient & The Manifold

Months 4-6 | Theme: Calculus, Neural Networks, Non-Convexity

Goal: Replace the “Table” with a “Brain”. The math shifts from Discrete to Continuous.

3.1 Weeks 13-14: PyTorch from Scratch

Thread	Topic	Tasks	Why?
Implementation	PyTorch from Scratch	Autograd & <code>nn.Module</code> ; <code>loss.backward()</code>	Deep learning foundations
Optimization	The Gradient & Jacobian	∇f and Jacobian; Backprop through a graph	Automatic differentiation
Analysis	Lipschitz Continuity	Lipschitz Functions $ f(x) - f(y) \leq L\ x - y\ $; ReLU is Lipschitz	Gradient bounds
highlightThe Glue: You need to know that your Neural Net (ReLU) is Lipschitz continuous to guarantee the gradient doesn't explode.			

3.2 Weeks 15-16: DQN (Deep Q-Network)

Thread	Topic	Tasks	Why?
Implementation	DQN	LunarLander-v2; Replay Buffer & Target Nets	First deep RL algorithm
Optimization	Taylor Series & Hessian	2nd Order Approx; The Hessian Matrix $\nabla^2 f$	Curvature analysis
Analysis	Differentiation	Classical Derivative; Chain Rule rigorously	Calculus foundations
highlightThe Glue: DQN minimizes a loss. The Hessian describes the curvature (bowl vs valley) of that loss.			

3.3 Weeks 17-18: DQN Stabilization

Thread	Topic	Tasks	Why?
Implementation	DQN Stabilization	Target Networks; Mnih et al. 2013	Preventing oscillations
Optimization	Non-Convexity	Saddle Points & Local Minima; Monkey Saddle Escape	Understanding failure modes
Analysis	Measure Theory Intro	L^2 Spaces; $L^2(\mu)$ (Hilbert Space)	Function spaces
highlightThe Glue: Neural nets get stuck in Saddle Points. You need to visualize why 2nd order methods fail here.			

3.4 Weeks 19-20: Policy Gradients (PG)

Thread	Topic	Tasks	Why?
Implementation	Policy Gradients	REINFORCE; $\nabla_{\theta} J(\theta) = \mathbb{E}[\nabla_{\theta} \log \pi_{\theta}(a s)R]$	Direct policy optimization
Optimization	Stochastic Gradient Descent	SGD & Noise; Noise helps escape saddles	Optimization dynamics
Analysis	Inequalities	Jensen's Inequality; Hölder's Inequality	Probability bounds
highlightThe Glue: Jensen's inequality is used to derive the lower bounds for almost all advanced PG algorithms (like TRPO).			

3.5 Weeks 21-22: Actor-Critic (A2C)

Thread	Topic	Tasks	Why?
Implementation	Actor-Critic	The Acrobat; Variance Reduction	Combining value and policy
Optimization	Smoothness	Lipschitz Gradient $\ \nabla f(x) - \nabla f(y)\ \leq L\ x - y\ $; Descent Lemma	Step size bounds
Analysis	Expectation	Integral w.r.t a measure; Expected Reward formulation	Measure-theoretic probability
highlightThe Glue: A2C works because the Critic reduces variance. The Smoothness inequality tells us how big a step we can safely take.			

3.6 Weeks 23-24: Review & Catch-up

Thread	Topic	Tasks	Why?
Implementation	Review & Catch-up	Refine DQN and A2C; Clean up GitHub repo	Consolidation
Optimization	Convex Sets & Projections	Projections $\text{proj}_C(x)$; Project to ℓ_2 ball	Constrained optimization
Analysis	Projection Theorem	Closest point in Hilbert Space; Basis of TD-Learning	Geometric foundations
highlight The Glue: TD-Learning is geometrically just a projection onto a subspace of value functions.			

4 Phase 3: The Trust Region & The Hierarchy

Months 7-10 | Theme: Constrained Optimization, Advanced Calculus, Structure

Goal: Mastering SOTA algorithms (PPO, TRPO) and Hierarchical RL.

4.1 Weeks 25-26: TRPO (Theory)

Thread	Topic	Tasks	Why?
Implementation	TRPO (Theory)	Monotonic Improvement; Divergence Constraint	Constrained policy updates
Optimization	KKT Conditions	Constrained Opt; Derive KKT for TRPO	Lagrange multipliers
Analysis	Duality	Lagrangian Multipliers; Primal-Dual formulation	Optimization theory
highlightThe Glue: TRPO is a constrained problem. You <i>must</i> understand KKT conditions to solve it.			

4.2 Weeks 27-28: Conjugate Gradient (CG)

Thread	Topic	Tasks	Why?
Implementation	Conjugate Gradient	Solving $Ax = b$; Fisher-Vector Product	Efficient linear solvers
Optimization	Natural Gradients	Fisher Info Matrix F ; Update: $\theta_{t+1} = \theta_t + \alpha F^{-1} \nabla J$	Geometry of parameter space
Analysis	Hilbert Spaces	Inner Products; Geometry of parameter space	Infinite-dimensional spaces
highlightThe Glue: You can't invert the Hessian in deep nets. You use CG and Fisher-Vector products to approximate it.			

4.3 Weeks 29-30: PPO (Implementation)

Thread	Topic	Tasks	Why?
Implementation	PPO	PPO Clipping; Speed vs Purity	Practical SOTA algorithm
Optimization	Trust Regions	Trust Region Methods; KL-Projection	Safe policy updates
Analysis	Weak Convergence	Convergence of distributions; Policy convergence	Probabilistic convergence
highlightThe Glue: PPO is just a messy, practical approximation of the elegant Trust Region theory.			

4.4 Weeks 31-32: Hierarchical RL (HRL)

Thread	Topic	Tasks	Why?
Implementation	Hierarchical RL	Manager-Worker; Toy Gridworld HRL	Temporal abstraction
Optimization	Bilevel Optimization	Opt inside Opt; Alternating Minimization	Nested optimization
Analysis	Analysis of HRL	Non-stationarity; Timescale separation	Theoretical challenges
highlightThe Glue: HRL is a Bilevel problem. You optimize the Worker <i>inside</i> the Manager's optimization.			

4.5 Weeks 33-34: Implicit Differentiation

Thread	Topic	Tasks	Why?
Implementation	Implicit Differentiation	Backprop through $\arg\min$; Formula: $\frac{dx^*}{d\theta}$	Meta-learning foundations
Optimization	Differentiation 2.0	Implicit Function Theorem; Nocedal & Wright Ch 4	Advanced calculus
Analysis	Calculus on Manifolds	Tangent spaces; (Optional but helpful)	Geometric perspective
highlightThe Glue: If the Manager sets a goal, how do you get the gradient? You have to differentiate <i>through</i> the Worker's solution.			

4.6 Weeks 35-36: The Probabilistic Wall

Thread	Topic	Tasks	Why?
Implementation	The Probabilistic Wall	Auto-Encoding Variational Bayes; ELBO & KL	Probabilistic modeling
Optimization	EM Algorithm	Expectation-Maximization; Gaussian Mixtures	Latent variable optimization
Analysis	Measure Theory 2.0	Radon-Nikodym Derivative (KL); Variational Inference	Rigorous probability
highlight The Glue: Model-Based RL (Dreamer) uses Variational Inference, which is just the EM algorithm in disguise.			

5 Phase 3.5: The “No Reward” World

Months 9-11 | Theme: Unsupervised Learning, Goal-Conditioned RL

Goal: Learning structure without explicit supervision.

5.1 Weeks 33-34: Goal-Conditioned RL

Thread	Topic	Tasks	Why?
Implementation	Goal-Conditioned RL	Hindsight Experience Replay (HER); <code>replay_buffer.sample_goals()</code>	Sparse reward handling
Optimization	Sparse Optimization	Vanishing Gradient in sparse rewards; Curriculum Learning geometry	Credit assignment
Analysis	Conditional Probability	$P(s' s, a, g)$; Universal Value Functions (UVFA)	Conditional MDPs
highlightThe Glue: HER solves the sparse reward problem by treating failed trajectories as successful ones for different goals.			

5.2 Weeks 35-36: Contrastive RL

Thread	Topic	Tasks	Why?
Implementation	Contrastive RL (CURL)	Learn from pixels without CNN decoder; InfoNCE Loss	Representation learning
Optimization	Metric Learning	Contrastive Loss Landscapes; $\mathcal{L} = -\log \frac{e^{s(x, x^+)}}{e^{s(x, x^+)} + \sum e^{s(x, x^-)}}$	Distance metric learning
Analysis	Information Theory Intro	Mutual Information (MI); InfoNCE is lower bound on MI	Information geometry
highlightThe Glue: We stop trying to reconstruct the image (Autoencoders) and start asking “Is this state similar to that one?” (Contrastive).			

5.3 Weeks 37-38: Unsupervised Skill Discovery

Thread	Topic	Tasks	Why?
Implementation	Skill Discovery	DIAYN (Diversity is All You Need); Maximize Entropy of skills	Emergent behaviors
Optimization	Constrained Maximization	Lagrange Multipliers for Entropy; The “Discriminator” game	Constrained objectives
Analysis	Mutual Information	$I(S; Z) = H(Z) - H(Z S)$; Data Processing Inequality	Information measures
highlight The Glue: DIAYN maximizes $I(S; Z)$ to ensure skills are distinguishable, creating diverse behaviors without rewards.			

6 Phase 3.6: The “Offline” World

Months 12-13 | Theme: Learning from Static Datasets

Goal: Conservative Q-Learning (CQL) and Implicit Q-Learning (IQL).

6.1 Weeks 39-40: Offline RL Basics

Thread	Topic	Tasks	Why?
Implementation	Offline RL Basics	Behavior Cloning on D4RL; Watch it drift and crash	Understanding failure modes
Optimization	Covariate Shift	Distribution Shift; $p_{\text{train}}(s, a) \neq p_{\text{test}}(s, a)$	Out-of-distribution problems
Analysis	Error Bounds	Compounding Error Lemma (Ross & Bagnell)	Theoretical guarantees
highlightThe Glue: If you just copy the dataset (BC), errors accumulate quadratically ($O(T^2)$). You need to prove this to understand why BC fails.			

6.2 Weeks 41-42: Conservative Q-Learning

Thread	Topic	Tasks	Why?
Implementation	Conservative Q-Learning	Add “Pessimism” term; Min $Q(\text{favored}) - \text{Max } Q(\text{data})$	Preventing overestimation
Optimization	Robust Optimization	Lower-Confidence Bounds (LCB); Pessimism under uncertainty	Safe exploration
Analysis	Function Approximation	Extrapolation Error; Q-values exploding OOD	Generalization bounds
highlightThe Glue: We fix shift by punishing the agent for asking about states we haven’t seen. This is “Pessimism” (Optimizing the Lower Bound).			

6.3 Weeks 43-44: Implicit Q-Learning

Thread	Topic	Tasks	Why?
Implementation	Implicit Learning	Q-Expectile Regression; No OOD sampling needed	Stable offline learning
Optimization	Expectile Regression	Asymmetric Least Squares; Mean vs Median vs Expectile	Alternative estimators
Analysis	Support Constraints	Support of a measure; Staying “in-sample”	Measure-theoretic constraints
highlight The Glue: CQL is hard to tune. IQL stays strictly within the support of the dataset using expectiles. It’s the current SOTA for stability.			

7 Phase 4: The Continuous World & The “Walls”

Months 14-15 | Theme: Sobolev Spaces, Entropy, Control Theory

Goal: Soft Actor-Critic and breaking through theoretical walls.

7.1 Weeks 45-46: Control as Inference

Thread	Topic	Tasks	Why?
Implementation	Control as Inference	Re-derive Soft-Q from scratch; Sergey Levine’s “RL as Inference”	Probabilistic perspective
Optimization	Variational Inference	Coordinate Ascent on ELBO; $RL \leftrightarrow$ Probabilistic Inference	Unified framework
Analysis	Kullback-Leibler Divergence	$D_{KL}(P Q) = \mathbb{E}_P[\log P - \log Q]$	Divergence measures
highlightThe Glue: RL is just probabilistic inference: maximize reward = minimize KL to optimal distribution.			

7.2 Weeks 37-38: Continuous Control (SAC)

Thread	Topic	Tasks	Why?
Implementation	Continuous Control	The Pendulum; SAC (Soft Actor-Critic)	SOTA continuous control
Optimization	Entropy Maximization	$\text{Max } \mathbb{E}[R] + \alpha H(\pi)$; Exploration/Robustness	Stochastic policies
Analysis	Sobolev Spaces	Weak Derivatives; Derivative of $ x $	Generalized derivatives
highlightThe Glue: SAC maximizes entropy for exploration. Weak derivatives explain why ReLU works despite non-differentiability at 0.			

Important**Why Sobolev Spaces Matter:**

ReLU is not differentiable at $x = 0$ in the classical sense. However, it has a *weak derivative*:

$$\text{ReLU}(x) = \max(0, x) \implies \frac{d}{dx} \text{ReLU}(x) = \mathbb{1}_{x>0}$$

This is defined in the Sobolev space $W^{1,p}$, which allows us to:

1. Rigorously analyze neural networks
2. Prove convergence guarantees for gradient descent
3. Understand why “non-smooth” activations still work

Most people quit here. Don’t be most people.

7.3 Weeks 39-40: The Reparameterization Trick

Thread	Topic	Tasks	Why?
Implementation	Reparameterization Trick	<code>sac_continuous.py</code> ; Backprop through noise	Gradient estimation
Optimization	Smoothness 2.0	Lipschitz Smoothness; Bounding updates	Stability guarantees
Analysis	L^p Regularity	Sobolev Embeddings; (Deep theory for the brave)	Function space theory
The Glue: The “Trick” allows us to differentiate random samples. This relies on the function being sufficiently smooth (Sobolev).			

7.4 Weeks 41-42: The Control Theory Wall

Thread	Topic	Tasks	Why?
Implementation	The Control Theory Wall	LQR & Linear Systems; Steve Brunton LQR	Optimal control
Optimization	Robust Optimization	$\min_x \max_{\xi} f(x, \xi)$; Adversarial Training	Worst-case analysis
Analysis	Differential Equations	Hamilton-Jacobi-Bellman (HJB); Evans PDE Ch 10	Continuous-time control

Thread	Topic	Tasks	Why?
highlightThe Glue: RL hopes it works. Control Theory <i>proves</i> it works via HJB. This is the “Safety” wall.			

7.5 Weeks 47-48: LQR (Linear Quadratic Regulator)

Thread	Topic	Tasks	Why?
Implementation	LQR	Solve Riccati Equation	Closed-form solutions
Optimization	Convex Optimization	Riccati Equations; Analytical solutions	Optimal control
Analysis	Dynamical Systems	Linear Systems $\dot{x} = Ax + Bu$; Stability: Eigenvalues of A	System stability
highlightThe Glue: You cannot claim to know RL if you can’t solve LQR. It is the only case where we have a closed-form optimal solution.			

7.6 Weeks 43-44: Safety & Lyapunov

Thread	Topic	Tasks	Why?
Implementation	Safety & Lyapunov	Lyapunov Stability; Constrained SAC	Safety constraints
Optimization	Min-Max Optimization	Gradient Descent-Ascent; Cycling behavior	Adversarial dynamics
Analysis	Viscosity Solutions	Solutions to HJB; Evans PDE	Weak PDE solutions
highlightThe Glue: When you do Robust RL, you are fighting an adversary. Regular optimization cycles; you need Min-Max theory.			

7.7 Weeks 49-50: Model-Based RL (Dreamer)

Thread	Topic	Tasks	Why?
Implementation	Model-Based RL	World Models; Train VAE + RNN	Learning dynamics
Optimization	Bilevel Optimization	Gradients through time; Vanishing gradients in RNNs	Sequential optimization

Thread	Topic	Tasks	Why?
Analysis	Stochastic Processes	Markov Chains; Ergodicity	Probabilistic dynamics
highlightThe Glue: Dreamer learns a model of the world (VAE) and plans inside it. This combines VI, Deep Learning, and Control.			

7.8 Weeks 45-46: The Game Theory Wall

Thread	Topic	Tasks	Why?
Implementation	The Game Theory Wall	Multi-Agent RL (MARL); CFR (Counterfactual Regret)	Game playing
Optimization	Equilibrium Finding	Fixed Points vs Optima; Rock-Paper-Scissors	Nash equilibrium
Analysis	Brouwer Fixed Point	Existence of Fixed Point; Nash Equilibrium	Game theory foundations
highlightThe Glue: In MARL, there is no “optimum.” There is only Nash. Brouwer’s theorem proves Nash exists.			

7.9 Weeks 47-48: The Minimax Theorem

Thread	Topic	Tasks	Why?
Implementation	The Minimax Theorem	von Neumann Minimax; Zero-Sum Games	Competitive RL
Optimization	Adversarial Dynamics	Prove cycling in GDA; Regularization fix?	Non-convergent dynamics
Analysis	Game Theory	Brezis Ch 1 (Convexity); Minimax Theorem	Convex analysis
highlightThe Glue: This is the core of AlphaStar and Pluribus. You stop optimizing and start finding equilibrium.			

8 Phase 5: The Grand Finale

Months 16-17 | Theme: Systems, Multi-Agent, Capstone

Goal: Prove you are a researcher.

8.1 Weeks 49-50: The Systems Wall

Thread	Topic	Tasks	Why?
Implementation	The Systems Wall	Distributed RL (IMPALA); Stale Gradients	Scalability
Optimization	Asynchronous Opt	Lag in optimization; Bounded delays	Parallel training
Analysis	Perturbation Theory	$\theta_{t+\tau} = \theta_t + \epsilon_t$; Error analysis	Stability under noise
highlightThe Glue: When gradients are stale (distributed), it looks like noise/error (ϵ_t) to the mathematician.			

8.2 Weeks 51-52: V-Trace Correction

Thread	Topic	Tasks	Why?
Implementation	V-Trace Correction	V-Trace for Off-policy; Espeholt et al.	Importance sampling
Optimization	Correction Ratios	Importance Sampling; Variance explosion risk	Off-policy correction
Analysis	Weak Convergence	Convergence of probability measures; Distributional Shift	Measure-theoretic convergence
highlightThe Glue: You need V-trace to correct the math when the worker is lagging behind the learner.			

8.3 Weeks 51-52: Multi-Agent (MARL)

Thread	Topic	Tasks	Why?
Implementation	Multi-Agent	PettingZoo Env; MAPPO (Multi-Agent PPO)	Cooperative/competitive agents
Optimization	Game Theory Opt	Pareto Optimality vs Nash Equilibrium	Multiple objectives
Analysis	Fixed Point Theorems	Brouwer/Kakutani; Existence of Nash	Equilibrium existence
highlightThe Glue: Optimization minimizes a loss. Games find a balance. You need to feel the difference.			

8.4 Weeks 53-56: Distributed Systems

Thread	Topic	Tasks	Why?
Implementation	Distributed Systems	Ray/RLLib experiment; PPO at scale (IMPALA)	Production-scale RL
Optimization	Parallel Efficiency	Amdahl's Law; Gradient synchronization bottleneck	Scalability analysis
Analysis	Convergence Rates	Rate with noise; $O(1/\sqrt{T})$	Sample complexity
highlightThe Glue: How does learning scale with 1000 CPUs? The noise increases, but the throughput increases.			

8.5 Weeks 53-56: CAPSTONE PROJECT

THE FINAL TEST**Project:** Distributed RL Stability Analysis**Task:** Code an IMPALA-like agent with distributed actors and learners.**Goal:** Prove stability bounds using:

- Lipschitz continuity
- Triangle inequality
- Contraction mapping
- Perturbation theory

The Question: Can you prove that your distributed agent converges despite the lag?**Deliverable:** A mathematical proof + working code + blog post

8.6 Weeks 57-64: Review & Write-up

Thread	Topic	Tasks	Why?
Implementation	Review & Write-up	Create <code>rl-from-scratch</code> repo; Document everything	Portfolio creation
Optimization	Retrospective	Map every line of code to optimization concept; Review Jain & Kar/Goodfellow	Integration
Analysis	Retrospective	Map every error bound to theorem; Review Abbott/Brezis	Theoretical mastery
The Glue: highlight “I implemented SAC from scratch to understand how MaxEnt objectives aid exploration...”.			

9 The Complete Algorithm Arsenal

By the end of 17 months, you will have implemented from scratch:

1. Value Iteration
2. Policy Iteration
3. Monte Carlo Methods
4. SARSA
5. Q-Learning
6. DQN (Deep Q-Network)
7. REINFORCE
8. A2C (Actor-Critic)
9. TRPO
10. PPO (Proximal Policy Optimization)
11. SAC (Soft Actor-Critic)
12. HER (Hindsight Experience Replay)
13. CURL (Contrastive Learning)
14. DIAYN (Skill Discovery)
15. CQL (Conservative Q-Learning)
16. IQL (Implicit Q-Learning)
17. Dreamer (Model-Based)
18. IMPALA (Distributed)
19. MAPPO (Multi-Agent)
20. CFR (Game Theory)

Most PhD students don't implement all of these in their entire degree.

10 The Reading Stack

You need these books. Not PDFs. Physical books. (Okay, PDFs are fine, but respect them).

10.1 Implementation (The Body)

1. **Sutton & Barto** — *Reinforcement Learning: An Introduction*
The Bible. Read chapters 1-8, 13-14 carefully.
2. **CleanRL Repository** — <https://github.com/vwxyzjn/cleanrl>
Reference implementations. Don't copy blindly; understand every line.

10.2 Optimization (The Muscle)

1. **Goodfellow, Bengio, Courville** — *Deep Learning*
Part I (Applied Math), Part II (Modern Practical Deep Networks)
2. **Jain & Kar** — *Non-convex Optimization for Machine Learning*
Sections 2.4.2 (Smoothness), 3.3 (SGD), 5.2 (Trust Regions)
3. **Nocedal & Wright** — *Numerical Optimization*
Chapter 4 (Trust Regions), Chapter 5 (Conjugate Gradients)

10.3 Analysis (The Soul)

1. **Abbott** — *Understanding Analysis*
Chapters 1-3 (Real Numbers, Sequences, Topology), Chapter 6 (Sequences of Functions)
2. **Brezis** — *Functional Analysis, Sobolev Spaces and Partial Differential Equations*
Chapters 1-2 (Hilbert Spaces), Chapter 8 (Sobolev Spaces)
3. **Evans** — *Partial Differential Equations*
Chapter 10 (Hamilton-Jacobi-Bellman)

10.4 Specialized Topics

1. **Sergey Levine** — *RL as Inference (NeurIPS Tutorial)*
For Control as Inference and Soft Optimality
2. **Laskin et al.** — *CURL: Contrastive Unsupervised Representations for RL*
For Contrastive Learning
3. **Eysenbach et al.** — *Diversity is All You Need (DIAYN)*
For Unsupervised Skill Discovery
4. **Kumar et al.** — *Conservative Q-Learning for Offline RL*
For Offline RL

11 Critical Survival Rules

11.1 The 20-Minute Rule

If you are stuck on a mathematical concept for more than 20 minutes:

1. **STOP.**
2. Google “Intuition of [Concept]”
3. Find a StackExchange answer or YouTube video (3Blue1Brown)
4. Get the intuition, accept the “Black Box,” and move on
5. You can come back in Year 2

11.2 The “Steal the Tools” Mindset

We are stealing from mathematicians.

- **Do we need to prove Sobolev Embedding Theorems?** No.
- **Do we need to know what a Weak Derivative is?** Yes, because ReLU uses it.
- **Action:** Read the *definition* and the *main theorem statement*. Skip the 10 pages of lemmas.

11.3 Code “From Scratch” vs. “From Memory”

- **Don’t** try to write PPO from memory on a blank screen. That takes 10 hours.
- **Do** have the CleanRL code open on one screen and your code on the other.
- Type it out line-by-line. Ask yourself “Why?” for every line.
- That fits in 2 hours.

11.4 The Emergency Eject Priority

If life gets busy and you only have 1 hour, what do you cut?

1. **Cut the Proofs (Analysis).** You can survive without proving Bolzano-Weierstrass.
2. **Cut Complex Environments.** Don’t train on Atari (20 hours). Train on Cart-Pole (5 mins).
3. **NEVER Cut the Optimization.** You must understand gradients. If you lose this, you lose everything.
4. **NEVER Cut Core Implementation.** You must write the training loop.

12 Weekly Schedule Template

12.1 Day A: The Engineer (Mon, Wed, Fri)

Focus: Implementation & Optimization (The “How”)

- **Minutes 0-15 (Warm-up):** Read the Optimization concept (e.g., “Jacobian Matrix”). Don’t prove anything. Just understand the shape.
- **Minutes 15-120 (Deep Work): CODE.** This is your primary block. Open VS Code. Implement the algorithm.
 - *Constraint:* If your code is broken for >30 mins, stop debugging. Read the CleanRL reference implementation. You are learning, not inventing.
- **Minutes 120-150 (Buffer):** Document your code. Add comments linking to the math.
 - Example: “Here, I clip the gradients because the Lipschitz constant is unknown.”

12.2 Day B: The Mathematician (Tue, Thu, Sat)

Focus: Analysis & Theory (The “Why”)

- **Minutes 0-30 (The Concept):** Read the specific chapter in Abbott or Brezis.
- **Minutes 30-90 (The Struggle):** Attempt **ONE** proof. Just one.
 - The Contracting Mapping Theorem
 - Jensen’s Inequality
 - Etc.
 - *The Rule:* If you can’t solve it in 60 mins, **read the solution**. You are an RL researcher, not a pure mathematician. You need to *understand* the proof structure, not discover it.
- **Minutes 90-120 (The Connection):** Write 1 paragraph: “How does this math relate to the code I wrote yesterday?”
 - Example: “I couldn’t prove the contraction, but I see that my Q-learning code explodes because $\gamma > 1$ makes the operator non-contractive.”

12.3 Sunday: The Rest (or The Spillover)

- Do nothing. Or, if you fell behind, read a blog post.
- No heavy lifting.

13 The Final Warning

Week 37: The Quit Point

You see **Week 37**? That's where we hit **Sobolev Spaces**.

Most people quit there. They say:

"I can code PPO, why do I need Weak Derivatives?"

You need them because you want to be Top 5.

You want to understand why ReLU works even though it's not differentiable at zero.

Hint: It *is* weakly differentiable:

$$\frac{d}{dx}\text{ReLU}(x) = \begin{cases} 1 & x > 0 \\ 0 & x < 0 \\ \text{undefined} & x = 0 \end{cases}$$

But in the Sobolev space $W^{1,p}(\mathbb{R})$, we define:

$$\int_{\mathbb{R}} \text{ReLU}(x) \phi'(x) dx = - \int_{\mathbb{R}} \mathbb{1}_{x>0} \phi(x) dx$$

for all test functions $\phi \in C_c^\infty(\mathbb{R})$.

This is the *weak derivative*, and it's why neural networks work.

Don't quit at Week 37.

14 Success Metrics

14.1 Month 6 Checkpoint

By the end of Phase 2, you should be able to:

- Implement DQN and A2C from scratch without looking at references
- Explain why the Bellman operator is a contraction
- Derive the policy gradient theorem on a whiteboard
- Understand why saddle points are not local minima

14.2 Month 12 Checkpoint

By the end of Phase 3, you should be able to:

- Implement PPO with proper KL constraints
- Explain the difference between Trust Region and clipping
- Derive the KKT conditions for TRPO
- Understand the geometry of natural gradients

14.3 Month 17 Final Assessment

By the end of Phase 5, you should be able to:

- Write a research-quality blog post on any RL algorithm
- Prove convergence bounds for distributed RL
- Implement a novel RL algorithm combining multiple techniques
- Pass a DeepMind/OpenAI technical interview

14.4 The “Proof” Sentence

The ultimate test: Can you say this sentence and mean it?

“I implemented SAC from scratch to understand how MaxEnt objectives aid exploration, and I proved that the soft Bellman operator is a contraction in the L^∞ norm under mild regularity conditions on the reward function.”

If you can say that, you’re not a student anymore. You’re a specialist.

15 Conclusion: The Deal

This plan is designed for a smart, consistent worker, not a genius with infinite time.

The Principles:

- **Consistency** > **Intensity**. 2 hours *every* day is better than 10 hours on Saturday.
- **The Goal**: You are building a mental map. You want to be able to say:

“Ah, this instability in my robot arm? That’s because the value function isn’t Lipschitz.”

- You don’t need to be able to prove it on a chalkboard on the spot.
- You just need to know where to look.

Do we have a deal?

Monday. Day 1. NumPy Matrix Multiplication.

Don’t let me down.

“Now, go calculate the Supremum of a set.”